

Anvil: executable verification of LLM-generated HPC operational artifacts

Antonio Rotundo

antonio.rotundo2@studio.unibo.it

August 6, 2026

Abstract

Assistants that write SLURM job scripts for supercomputer users are evaluated with similarity metrics, because no benchmark verifies whether the artifacts they produce actually work. Anvil verifies them by execution against a declared reference cluster: five levels covering syntax, submittability, execution, effective resource fit and safety, bracketed by an oracle that must score 1.0 and a deliberately faulty model that must score 0.0. On eight from-scratch tasks and 220 induced repair tasks we measure three models across two families, and report four results, two of them deflationary about the benchmark’s own machinery. Submittability is not ordered by model size: a 3B model of one family beats a 7B of another, and within one family the level falls as size rises. Real job submission changes exactly one verdict in 3204 artifacts against a **bash** sandbox, so the most expensive component of the harness is nearly redundant on this task set. Retrieved documentation costs the smaller model 28 points of resource fit whatever the text says, while the larger one is unaffected by irrelevant text and loses 6 points to the one document that addresses the level it concerns. A task built so that GNU coreutils and **utils** must judge it differently is failed by all three models for reasons that precede the difference. We also report, rather than conceal, three measurement faults found and corrected during the work, one of which had already been published.

1 Introduction

A wrong `--mem` does not look wrong. It looks plausible, and the job dies at hour six.

Operational HPC artifacts, batch scripts and container recipes, sit in a gap. Execution-based evaluation is standard for program synthesis and increasingly common for infrastructure code, yet the assistants appearing for HPC user support are still scored by how closely their output resembles a reference answer. Resemblance is the wrong instrument here for a specific reason: the failures that matter are invisible in the text. A directive placed after the first command is silently ignored by the scheduler. An omitted `--ntasks` is not an absent request but a request for a different number. A memory figure one factor too small parses, submits, and runs until it does not.

Unlike cloud infrastructure code, where a plan cannot be executed cheaply and evaluation falls back on graph edit distance or model judges, these artifacts are directly checkable. `sbatch --test-only` costs milliseconds, the job really runs, and the ground truth is objective. The empty slot is not a hard problem; it is an unoccupied one.

This paper reports what the resulting instrument says about three models and, with more care, what it says about itself.

2 Related work

This section is not written. The repository ships `scripts/litsweep.py`, a reproducible sweep over OpenAlex, Crossref, DBLP and arXiv, so that the review can be regenerated and cited rather than asserted. Running it and citing its output is the remaining work on this manuscript. No references are given here rather than plausible ones: a citation nobody can follow would compromise the only thing this paper offers, which is that its claims can be checked.

The single prior result the method depends on directly is the unbiased $\text{pass}@k$ estimator of Chen et al. [1], used unchanged.

3 The benchmark

T1, writing from scratch. Eight tasks, each a support request in ordinary words paired with a machine-readable specification: node and task counts, walltime and memory bounds, GPUs, directives the task demands explicitly, and strings the script must print when it runs.

T2, diagnosing and repairing. 220 tasks derived mechanically from the T1 canonical solutions by seven fault injectors, anchored to classes observed on a real model rather than invented: an omitted directive whose SLURM default hides it, a misplaced directive, prose leaking into a value, an absent no-default directive, a script stripped of every directive, a payload that no longer matches its own specification, and a value the scheduler rejects. Induction keeps only the variants that actually fail verification, since an injector producing an accidentally valid script is a bug in the injector and not a fault worth teaching a model to repair. Repairs are graded by the same verifier as T1, with no partial credit for being closer to correct.

4 Method

Five levels. The strict metric requires all of them.

syntax a shebang, `bash -n`, and no directive after the first command.

submittability `sbatch --test-only` accepts it.

functional it runs and exits zero.

resource fit the *effective* request matches the specification, with SLURM’s documented defaults applied.

safety destructive-pattern probes, which gate execution.

A skipped level is never a passed one. Without a scheduler, submittability is skipped and scored as not passed. One refinement was forced by measurement: a level skipped because *this machine* lacks a capability is pass-through, since every artifact is skipped alike, whereas a level skipped because *this artifact* cannot be judged anywhere is a failure. Before the distinction existed, ten repairs of a dependency task passed strict by being unjudgeable.

Effective requests, not string presence. A serial script omitting `--nodes` still requests one node and is correct. An early version checked for the string and failed scripts the scheduler accepts, which is precisely the surface-form matching this work exists to replace.

Both bounds shipped. An oracle returning canonical solutions must score 1.0 or the benchmark is defective; a broken model returning faulty artifacts must score 0.0 strict, and its sampling is arranged so the destructive flavour is always drawn, because a guard that never fires is decoration.

Table 1: T1, writing a job script from scratch. pass@1 per level. Safety is 1.000 ± 0.000 for every model and is omitted.

model	syntax	submit.	functional	res. fit	strict
Qwen2.5-Coder 1.5B	0.575 ± 0.025	0.842 ± 0.013	0.533 ± 0.037	0.442 ± 0.013	0.308 ± 0.025
Granite 4.1 3B	1.000 ± 0.000	0.875 ± 0.000	0.842 ± 0.037	0.625 ± 0.000	0.500 ± 0.000
Qwen2.5-Coder 7B	1.000 ± 0.000	0.792 ± 0.025	0.875 ± 0.000	1.000 ± 0.000	0.667 ± 0.025

Table 2: T2, diagnosing and repairing an induced fault. pass@1 per level. Safety is 1.000 ± 0.000 for every model and is omitted.

model	syntax	submit.	functional	res. fit	strict
Qwen2.5-Coder 1.5B	0.792 ± 0.016	0.886 ± 0.005	0.664 ± 0.005	0.412 ± 0.014	0.291 ± 0.000
Granite 4.1 3B	0.862 ± 0.002	1.000 ± 0.000	0.750 ± 0.000	0.753 ± 0.009	0.661 ± 0.009
Qwen2.5-Coder 7B	0.983 ± 0.002	0.977 ± 0.000	0.870 ± 0.002	0.965 ± 0.007	0.824 ± 0.002

A declared reference cluster. Four nodes, sixteen cores and one GPU each, in a container running a live `slurmctld`, so that scores do not depend on the hardware of whoever runs the benchmark. Two preflights protect it: a canary that distinguishes a failing scheduler from a failing model, and a topology canary that refuses to score submittability when the scheduler in front of it is not the declared cluster.

Generation and verification are separate. Generation needs an accelerator; faithful grading needs the scheduler and GNU coreutils. Generations travel as JSONL carrying the digest of the task file, and the verifier refuses any whose digest does not match.

5 Results

Three seeds, $n = 5$ samples per task, 4-bit quantization on an RTX 3060, graded inside the container under the `bash` executor. The $\pm$ is half the range across seeds. It is not a confidence interval, and no significance is claimed anywhere in this paper.

5.1 Submittability is not ordered by size

On T1 the level reads 0.875 for a 3B model, 0.842 for a 1.5B and 0.792 for a 7B; on T2, 1.000, 0.977 and 0.886. The smallest model of the second family leads both, and within the Qwen family the level falls as size rises while every other level improves, which Figure 1 shows as the one bar group running against the others.

The two families also fail it differently. One invents partition names, which a different cluster might genuinely have; the other invents an option SLURM does not have, which no cluster has. The level measures how much a model volunteers that nothing asked for, and whose vocabulary it reaches for when it does.

5.2 Real submission is nearly redundant here

Grading the same generations twice inside the same image, once through a `bash` sandbox and once through real `sbatch`, exactly one artifact of 3204 changes verdict, and it changes in favour of real submission, which accepts a script the sandbox wrongly rejects. Every script real submission stopped was already failing another level.

This is a statement about these eight tasks and not about execution. A task built to need it does need it: a memory under-request that only cgroup enforcement can observe is invisible




Figure 1: T1 per level. Every level rises with capability except submittability, where the 3B model leads and the 7B trails both others.




Figure 2: Resource fit under four retrieval conditions. The small model pays a toll for any attached text, including a length-matched passage on an unrelated subject; the large model pays none, and loses ground only to the document about the two directives this level checks.

to every static level and to the sandbox, and there the verdict does change.

5.3 Retrieval costs, and the sign depends on prior knowledge

Seven conditions across two sizes, four of them shared and plotted in Figure 2. Any attached text costs the 1.5B about 28 points of resource fit, whether it is SLURM documentation or a passage about coastal tides cut to the same length. The 7B pays no such toll: two unrelated off-topic controls leave it on its zero-shot values to the digit. The single corpus document addressing `--time` and `--mem` is worth +21 points to the model that does not know those directives have no default, and -6 to the one that does.

Capacity decides whether a model can ignore what does not concern it, and relevance decides the sign of what it cannot ignore.

5.4 A toolchain difference no model reaches

GNU `coreutils` 9.4 and `utils` 0.8.0 agree on 91 of 101 probed invocations. Three of the disagreements are behavioural and need no misconfigured locale: character counting and column expansion on a non-ASCII payload, and SI number formatting. A task built to sit in that corner

is judged differently by the two implementations, which a guard asserts.

Asked to solve it, three models produced 45 samples: none pinned a locale, and none diverged. Each fails on both implementations for a reason that precedes the difference, one of them by emitting escape sequences as literal text and so removing the non-ASCII content the question was about.

6 What went wrong while measuring

Reported because the corrections are load-bearing.

A published multi-seed table had been graded against the experiment machine’s own scheduler, which has one node, no GPUs, and rejects three of the eight canonical solutions. Submitability and strict were wrong by up to 36 points, while syntax, functional and resource fit were unchanged to the digit, which is the signature that identified the fault. The topology canary exists because of it, and fired on that same machine the first time a model was run after it landed.

The same defect was later found in the retrieval ablation, whose two scheduler-dependent columns were withdrawn and remeasured. Correcting them inverted an arm ordering: a single-seed pilot had been right, and the larger run that appeared to correct it had been wrong. The cause was the grading environment, not the seeds.

Three mechanisms proposed for the retrieval effect were each refuted by the experiment run to test them. The fourth was found only by counting per sample instead of per problem.

7 Threats to validity

Eight T1 tasks, which makes each one worth 0.125 of every T1 figure. Three seeds, so the ranges are spread rather than intervals. Two model families at three sizes, all quantized to 4 bit, none large. One reference topology, declared by this work rather than borrowed from a centre in production. All prompts written by one author, in one language and one style. No result here has been replicated independently.

8 Availability

Code, tasks, canonical solutions, the digest manifest identifying the dataset version, the generated leaderboard and every figure above are at <https://github.com/antoniorotundo2/anvil> under MIT. The command `anvil check` applies the verifier to a single artifact, with or without a site policy, for readers who want the instrument rather than the measurements.

References

- [1] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.