

Anvil: executable verification of LLM-generated HPC operational artifacts

Antonio Rotundo

antonio.rotundo2@studio.unibo.it

August 9, 2026

Abstract

Assistants that write SLURM job scripts for supercomputer users are evaluated with similarity metrics, because no benchmark verifies whether the artifacts they produce actually work. Anvil verifies them by execution against a declared reference cluster: five levels covering syntax, submittability, execution, effective resource fit and safety, bracketed by an oracle that must score 1.0 and a deliberately faulty model that must score 0.0. On eight from-scratch tasks and 220 induced repair tasks we measure five models across three families and two generations of one of them. Submittability is not ordered by model size: on the from-scratch set a 3B model of one family and a 12B of another lead it while a 7B trails both, and within one family the level falls as size rises. Real job submission looked nearly redundant against a `bash` sandbox, one verdict in 3120 artifacts, until a fifth model wrote `srun` inside its scripts: the sandbox has no allocation for a job step to attach to, and the count moved to six in 3900 with 32 samples failing in the sandbox and passing on the scheduler. On a task set whose memory need only the payload knows, the same comparison reads 297 of 900. Retrieved documentation costs the smaller model 28 points of resource fit whatever the text says, while the larger one is unaffected by irrelevant text and loses 6 points to the one document that addresses the level it concerns. A task built so that GNU coreutils and `utils` must judge it differently is failed by all three models asked for reasons that precede the difference. We also report, rather than conceal, three measurement faults found and corrected during the work, one of which had already been published.

1 Introduction

A wrong `--mem` does not look wrong. It looks plausible, and the job dies at hour six.

Operational HPC artifacts, batch scripts and container recipes, sit in a gap. Execution-based evaluation is standard for program synthesis and increasingly common for infrastructure code, yet the assistants appearing for HPC user support are still scored by how closely their output resembles a reference answer. Resemblance is the wrong instrument here for a specific reason: the failures that matter are invisible in the text. A directive placed after the first command is silently ignored by the scheduler. An omitted `--ntasks` is not an absent request but a request for a different number. A memory figure one factor too small parses, submits, and runs until it does not.

Unlike cloud infrastructure code, where a plan cannot be executed cheaply and evaluation falls back on graph edit distance or model judges, these artifacts are directly checkable. `sbatch --test-only` costs milliseconds, the job really runs, and the ground truth is objective. The empty slot is not a hard problem; it is an unoccupied one.

This paper reports what the resulting instrument says about five models and, with more care, what it says about itself, including one claim that four models supported and a fifth overturned.

2 Related work

The works below were selected from a sweep of fourteen queries against OpenAlex, Crossref, DBLP and arXiv, run by `scripts/litsweep.py` and returning 761 unique records. Its full output is committed with this manuscript, so the selection can be audited rather than taken on trust.

Assistants for HPC users, and how they are judged. Support assistants for super-computer users are being built and deployed: a chatbot for autonomous user support [6], a science-gateway assistant whose outcomes were reported from a national laboratory [1], and a Slurm-native serving infrastructure for running such models on the cluster itself [3]. Zheng et al. [10] compare several of these approaches. The task predates the models: the NERSC job script generator [8] produced batch scripts from a form, with correctness guaranteed by construction rather than measured. What is missing across this line is an execution-grounded way to say whether the artifacts these systems emit would run, which is the gap this work addresses.

Language models for HPC code. A parallel line targets the source code rather than the operational artifacts around it: modelling parallel programs [5], zero-shot parallelization guided by graph neural networks [4], and MPI code generation for HPC workloads [9]. These evaluate a kernel that compiles and runs. A batch script is judged by a scheduler instead, and fails in ways a compiler has no opinion about: a directive after the first command is accepted and ignored, and an omitted one is not an absent request but a different one.

Execution-grounded evaluation elsewhere. Outside HPC, execution is the norm. Chen et al. [2] established the unbiased $\text{pass}@k$ estimator this work uses unchanged. Infrastructure-as-code is the closest neighbour and shows the constraint: unable to execute a plan cheaply, Santos Vargas et al. [7] benchmark generated Terraform through static analysis. HPC operational artifacts do not share that constraint. `sbatch --test-only` costs milliseconds and the job really runs, which is why the ground truth here can be objective where theirs cannot.

3 The benchmark

T1, writing from scratch. Eight tasks, each a support request in ordinary words paired with a machine-readable specification: node and task counts, walltime and memory bounds, GPUs, directives the task demands explicitly, and strings the script must print when it runs.

T2, diagnosing and repairing. 220 tasks derived mechanically from the T1 canonical solutions by seven fault injectors, anchored to classes observed on a real model rather than invented: an omitted directive whose SLURM default hides it, a misplaced directive, prose leaking into a value, an absent no-default directive, a script stripped of every directive, a payload that no longer matches its own specification, and a value the scheduler rejects. Induction keeps only the variants that actually fail verification, since an injector producing an accidentally valid script is a bug in the injector and not a fault worth teaching a model to repair. Repairs are graded by the same verifier as T1, with no partial credit for being closer to correct.

4 Method

Five levels. The strict metric requires all of them.

syntax a shebang, `bash -n`, and no directive after the first command.

submittability sbatch --test-only accepts it.

functional it runs and exits zero.

resource fit the *effective* request matches the specification, with SLURM’s documented defaults applied.

safety destructive-pattern probes, which gate execution.

A skipped level is never a passed one. Without a scheduler, submittability is skipped and scored as not passed. One refinement was forced by measurement: a level skipped because *this machine* lacks a capability is pass-through, since every artifact is skipped alike, whereas a level skipped because *this artifact* cannot be judged anywhere is a failure. Before the distinction existed, ten repairs of a dependency task passed strict by being unjudgeable.

Effective requests, not string presence. A serial script omitting --nodes still requests one node and is correct. An early version checked for the string and failed scripts the scheduler accepts, which is precisely the surface-form matching this work exists to replace.

Both bounds shipped. An oracle returning canonical solutions must score 1.0 or the benchmark is defective; a broken model returning faulty artifacts must score 0.0 strict, and its sampling is arranged so the destructive flavour is always drawn, because a guard that never fires is decoration.

A declared reference cluster. Four nodes, sixteen cores and one GPU each, in a container running a live `slurmctld`, so that scores do not depend on the hardware of whoever runs the benchmark. Two preflights protect it: a canary that distinguishes a failing scheduler from a failing model, and a topology canary that refuses to score submittability when the scheduler in front of it is not the declared cluster.

Generation and verification are separate. Generation needs an accelerator; faithful grading needs the scheduler and GNU coreutils. Generations travel as JSONL carrying the digest of the task file, and the verifier refuses any whose digest does not match.

5 Results

Three seeds, $n = 5$ samples per task, 4-bit quantization on an RTX 3060, graded inside the container under the `bash` executor. The $\pm$ is half the range across seeds. It is not a confidence interval, and no significance is claimed anywhere in this paper.

5.1 Submittability is not ordered by size

On T1 the level reads 0.875 for a 3B model and for a 9B of a third generation, 0.867 for a 12B, 0.842 for a 1.5B and 0.792 for a 7B. Four models sit within eight points of each other at the top and the 7B is alone at the bottom, while within the Qwen family the level falls as size rises and every other level improves, which Figure 1 shows as the one bar group running against the others. Two families of three are above Qwen across a range from 3B to 12B, so the shortfall is not one model’s quirk.

T2 does not repeat that ordering, and it would be overclaiming to say otherwise: there every model sits between 0.876 and 1.000, which leaves little room to separate them, and the family that leads T1 is last.

Table 1: T1, writing a job script from scratch. pass@1 per level. A row is a condition and not a model: the executor decides how *functional* is judged and the quantization decides how the model was loaded, and both are measurements rather than settings, so each is published rather than one being chosen. Only the four other levels differ between the **bash** and **sbatch** rows of one model. Safety is 1.000 ± 0.000 throughout and is omitted.

model, executor and loading	syntax	submit.	functional	res. fit	strict
Qwen2.5-Coder 1.5B (bash, 4-bit)	0.575 ± 0.025	0.842 ± 0.013	0.533 ± 0.037	0.442 ± 0.013	0.308 ± 0.025
Qwen2.5-Coder 1.5B (sbatch, 4-bit)	0.575 ± 0.025	0.842 ± 0.013	0.375 ± 0.025	0.442 ± 0.013	0.308 ± 0.025
Qwen2.5-Coder 1.5B (bash, fp16)	0.750 ± 0.000	1.000 ± 0.000	0.533 ± 0.025	0.617 ± 0.013	0.408 ± 0.025
Qwen3.5 9B (bash, 4-bit)	1.000 ± 0.000	0.875 ± 0.025	0.650 ± 0.025	0.600 ± 0.050	0.450 ± 0.025
Qwen3.5 9B (sbatch, 4-bit)	1.000 ± 0.000	0.875 ± 0.025	0.658 ± 0.013	0.600 ± 0.050	0.475 ± 0.025
Granite 4.1 3B (bash, 4-bit)	1.000 ± 0.000	0.875 ± 0.000	0.842 ± 0.037	0.550 ± 0.000	0.425 ± 0.000
Granite 4.1 3B (sbatch, 4-bit)	1.000 ± 0.000	0.875 ± 0.000	0.717 ± 0.037	0.550 ± 0.000	0.425 ± 0.000
Gemma 4 12B (bash, 4-bit)	0.875 ± 0.000	0.867 ± 0.013	0.875 ± 0.000	0.917 ± 0.050	0.658 ± 0.062
Gemma 4 12B (sbatch, 4-bit)	0.875 ± 0.000	0.867 ± 0.013	0.742 ± 0.013	0.917 ± 0.050	0.658 ± 0.062
Qwen2.5-Coder 7B (bash, 4-bit)	1.000 ± 0.000	0.792 ± 0.025	0.875 ± 0.000	1.000 ± 0.000	0.667 ± 0.025
Qwen2.5-Coder 7B (sbatch, 4-bit)	1.000 ± 0.000	0.792 ± 0.025	0.667 ± 0.025	1.000 ± 0.000	0.667 ± 0.025

Table 2: T2, diagnosing and repairing an induced fault. pass@1 per level, one row per condition as in Table 1. Safety is 1.000 ± 0.000 throughout and is omitted.

model, executor and loading	syntax	submit.	functional	res. fit	strict
Qwen2.5-Coder 1.5B (bash)	0.792 ± 0.016	0.886 ± 0.005	0.664 ± 0.005	0.330 ± 0.009	0.211 ± 0.007
Qwen2.5-Coder 1.5B (sbatch)	0.792 ± 0.016	0.886 ± 0.005	0.595 ± 0.009	0.330 ± 0.009	0.212 ± 0.005
Qwen3.5 9B (bash)	0.982 ± 0.005	0.982 ± 0.005	0.947 ± 0.011	0.711 ± 0.009	0.691 ± 0.009
Qwen3.5 9B (sbatch)	0.982 ± 0.005	0.982 ± 0.005	0.950 ± 0.009	0.711 ± 0.009	0.694 ± 0.009
Granite 4.1 3B (bash)	0.862 ± 0.002	1.000 ± 0.000	0.750 ± 0.000	0.621 ± 0.016	0.529 ± 0.016
Granite 4.1 3B (sbatch)	0.862 ± 0.002	1.000 ± 0.000	0.750 ± 0.000	0.621 ± 0.016	0.529 ± 0.016
Gemma 4 12B (bash)	1.000 ± 0.000	0.876 ± 0.002	0.885 ± 0.002	0.932 ± 0.005	0.726 ± 0.002
Gemma 4 12B (sbatch)	1.000 ± 0.000	0.876 ± 0.002	0.762 ± 0.002	0.932 ± 0.005	0.726 ± 0.002
Qwen2.5-Coder 7B (bash)	0.983 ± 0.002	0.977 ± 0.000	0.870 ± 0.002	0.965 ± 0.007	0.824 ± 0.002
Qwen2.5-Coder 7B (sbatch)	0.983 ± 0.002	0.977 ± 0.000	0.847 ± 0.002	0.965 ± 0.007	0.824 ± 0.002

The families also fail the level differently. One invents partition names, which a different cluster might genuinely have; another invents an option SLURM does not have, which no cluster has. The level measures how much a model volunteers that nothing asked for, and whose vocabulary it reaches for when it does.

5.2 Real submission was nearly redundant until a model wrote **srun**

Grading the same generations twice inside the same image, once through a **bash** sandbox and once through real **sbatch**, four models produced one changed verdict in 3120 artifacts. Everything real submission stopped had already failed another level, so the executor propagated a verdict rather than producing one, and this section said so.

The fifth model ended that. Six artifacts of 3900 now change verdict and 32 samples fail under the sandbox while passing under real submission, and five of the six share one cause:

```
srun: error: Unable to confirm allocation for job 1: Invalid job id specified
```

The model writes **srun** inside the batch script, which is how a job step is launched on a real cluster. The sandbox has no allocation for it to attach to, so the step exits non-zero and the sample is recorded as failing. The sandbox produces a false negative on every script that launches a job step, and did so silently while no model wrote one.

Both artifacts the two executors disagree about run the same direction, and it is not the one a sandbox is usually suspected of: neither is the scheduler catching something the sandbox




Figure 1: T1 per level, five models. Every level broadly rises with capability except submittability, where the 3B, the 9B and the 12B lead and the 7B trails all four.

Table 3: The execution-sensitive set from scratch, each model graded twice, pass@1 per level. Two tasks, three seeds, $n = 5$. The set states no memory minimum, so what a script needs is a property of the payload the model wrote and only running the job can decide it. Safety is 1.000 ± 0.000 throughout and is omitted.

model, executor and loading	syntax	submit.	functional	res. fit	strict
Qwen2.5-Coder 1.5B (bash)	1.000 ± 0.000	0.500 ± 0.200	0.333 ± 0.150	0.667 ± 0.050	0.133 ± 0.050
Qwen2.5-Coder 1.5B (sbatch)	1.000 ± 0.000	0.500 ± 0.200	0.167 ± 0.100	0.667 ± 0.050	0.133 ± 0.050
Qwen3.5 9B (bash)	0.833 ± 0.100	0.833 ± 0.100	0.467 ± 0.050	1.000 ± 0.000	0.467 ± 0.050
Qwen3.5 9B (sbatch)	0.833 ± 0.100	0.833 ± 0.100	0.700 ± 0.100	1.000 ± 0.000	0.700 ± 0.100
Granite 4.1 3B (bash)	1.000 ± 0.000	1.000 ± 0.000	0.467 ± 0.200	0.133 ± 0.050	0.067 ± 0.100
Granite 4.1 3B (sbatch)	1.000 ± 0.000	1.000 ± 0.000	0.467 ± 0.200	0.133 ± 0.050	0.067 ± 0.100
Gemma 4 12B (bash)	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000
Gemma 4 12B (sbatch)	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000	1.000 ± 0.000
Qwen2.5-Coder 7B (sbatch)	1.000 ± 0.000	0.500 ± 0.000	0.000 ± 0.000	0.500 ± 0.000	0.000 ± 0.000
Qwen2.5-Coder 7B (bash)	1.000 ± 0.000	0.500 ± 0.000	0.967 ± 0.050	0.500 ± 0.000	0.467 ± 0.050

missed, both are the sandbox rejecting what the scheduler accepts, for want of an allocation or a core binding. The deflationary claim was true of the four models measured and not of the method, which is the more useful way to have been wrong.

What it was really a claim about is the task set. Every task in the two sets above states its resource requirements in its prompt, so a static check can reach them and the executor has little left to decide. A second set states no memory minimum: the payload is written by the model, what it needs is a property of that payload, and no level that reads the text can know it. Graded the same way, five models and three seeds and 900 sample comparisons, the same executor comparison reads **297 artifacts of 900**, with 356 of the 394 stopped artifacts returning `OUT_OF_MEMORY`. From 0.15% to 33%, on the same harness and the same models.

Table 3 gives the from-scratch half, one row per arm. Qwen2.5-Coder 7B is the case: the sandbox promotes 29 of its 30 artifacts and the scheduler kills every one, so *functional* falls from 0.967 to 0.000 between its two rows. It is the model with 1.000 resource fit on the main T1 set, strong wherever the requirement is written down and at floor wherever the requirement is a property of what it wrote itself. Gemma 4 12B is the only model that solves the set, 1.000 under both arms.

Qwen3.5 9B moves the other way for the reason above, `srun` in the payload. On the repair half every model loses at least a third of its strict score under enforcement while the ordering is preserved, so this is a level of difficulty the sandbox cannot see rather than a re-ranking. The set




Figure 2: Resource fit under four retrieval conditions. The small model pays a toll for any attached text, including a length-matched passage on an unrelated subject; the large model pays none, and loses ground only to the document about the two directives this level checks.

is small, two tasks and ten repairs, and its numbers are not comparable with the 3900-sample tables; what it establishes is not a score but that the executor’s value is a property of what is being asked.

5.3 Retrieval costs, and the sign depends on prior knowledge

Seven conditions across two sizes, four of them shared and plotted in Figure 2. Any attached text costs the 1.5B about 28 points of resource fit, whether it is SLURM documentation or a passage about coastal tides cut to the same length. The 7B pays no such toll: two unrelated off-topic controls leave it on its zero-shot values to the digit. The single corpus document addressing `--time` and `--mem` is worth +21 points to the model that does not know those directives have no default, and −6 to the one that does.

Capacity decides whether a model can ignore what does not concern it, and relevance decides the sign of what it cannot ignore.

5.4 A toolchain difference no model reaches

GNU `coreutils` 9.4 and `utils` 0.8.0 agree on 91 of 101 probed invocations. Three of the disagreements are behavioural and need no misconfigured locale: character counting and column expansion on a non-ASCII payload, and SI number formatting. A task built to sit in that corner is judged differently by the two implementations, which a guard asserts.

Asked to solve it, three models produced 45 samples: none pinned a locale, and none diverged. Each fails on both implementations for a reason that precedes the difference, one of them by emitting escape sequences as literal text and so removing the non-ASCII content the question was about.

6 What went wrong while measuring

Reported because the corrections are load-bearing.

A published multi-seed table had been graded against the experiment machine’s own scheduler, which has one node, no GPUs, and rejects three of the eight canonical solutions. Submitability and strict were wrong by up to 36 points, while syntax, functional and resource fit were unchanged to the digit, which is the signature that identified the fault. The topology canary exists because of it, and fired on that same machine the first time a model was run after it landed.

The same defect was later found in the retrieval ablation, whose two scheduler-dependent columns were withdrawn and remeasured. Correcting them inverted an arm ordering: a single-seed pilot had been right, and the larger run that appeared to correct it had been wrong. The cause was the grading environment, not the seeds.

Three mechanisms proposed for the retrieval effect were each refuted by the experiment run to test them. The fourth was found only by counting per sample instead of per problem.

The resource fit check compared the requested walltime against the one a task names from above only, so a script asking fifteen seconds where the prompt asked fifteen minutes was scored correct. 123 of 2421 passing artifacts were requests of that kind. Regrading the saved generations under a floor moved five of ten published cells, left the other five identical to the digit, and dropped one model below another on strict: a one-sided comparison does not inflate a table evenly, it inflates whichever model has the habit the check cannot see. Reports and leaderboard entries now carry a digest of the verifier beside the digest of the task file, since a changed check invalidates a comparison exactly as thoroughly as a changed task and left no trace of having done so.

That gap surfaced while models were being screened for something else, so the remaining constraints were then audited on purpose. Memory had the same shape, compared from below only, and 30 of 2298 passes requested more than the task named, all of them repairs in which the model rewrote a directive the induced fault had not touched. Tightening it moved one cell, by the amount the audit had predicted to three decimals before the change was made. Requested GPU counts were exact in all 343 cases and needed nothing.

7 Threats to validity

Eight T1 tasks, which makes each one worth 0.125 of every T1 figure. Three seeds, so the ranges are spread rather than intervals. Three model families at five sizes, from 1.5B to 12B, all quantized to 4 bit; one of the five reasons by default and was measured with that disabled, since under this token budget it would otherwise be cut off before producing a script. One reference topology, declared by this work rather than borrowed from a centre in production. All prompts written by one author, in one language and one style. No result here has been replicated independently.

8 Availability

Code, tasks, canonical solutions, the digest manifest identifying the dataset version, the generated leaderboard and every figure above are at <https://github.com/antoniorotundo2/anvil> under MIT. The command `anvil check` applies the verifier to a single artifact, with or without a site policy, for readers who want the instrument rather than the measurements.

References

- [1] Brandon S. Biggs and Kaylee Dalton. Outcomes of HPC user support using a science gateway AI assistant. Technical report, Idaho National Laboratory, 2024.
- [2] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. doi: 10.48550/arXiv.2107.03374.

- [3] Ali Doosthosseini, Jonathan Decker, Hendrik Nolte, and Julian Kunkel. Chat AI: A seamless slurm-native solution for HPC-based services. *arXiv preprint arXiv:2407.00110*, 2024. doi: 10.48550/arXiv.2407.00110.
- [4] Quazi Ishtiaque Mahmud, Ali TehraniJamsaz, Hung Phan, Le Chen, and Mihai Capota. AutoParLLM: GNN-guided context generation for zero-shot code parallelization using LLMs. *arXiv preprint arXiv:2310.04047*, 2023. doi: 10.48550/arXiv.2310.04047.
- [5] Daniel Nichols, Aniruddha Marathe, Harshitha Menon, Todd Gamblin, and Abhinav Bhatele. HPC-coder: Modeling parallel programs using large language models. *arXiv preprint arXiv:2306.17281*, 2023. doi: 10.48550/arXiv.2306.17281.
- [6] Rajesh Pasupuleti, Ravi Vadapalli, Christopher Mader, and Subhan Shakoore. LLM chatbot for autonomous user support in high-performance computing. In *International Conference on Machine Learning and Applications (ICMLA)*, 2025. doi: 10.1109/icmla66185.2025.00226.
- [7] Francis Luis Santos Vargas, Rodrigo Brandão Mansilha, and Diego Kreutz. Can language models generate secure Terraform code? a security-focused benchmark using static analysis. In *Workshop on Cloud Computing Research*, 2026. doi: 10.5753/pesquisanuvem.2026.23954.
- [8] Robin Shao, Thorsten Kruth, and Zhengji Zhao. NERSC job script generator. In *IEEE/ACM International Workshop on HPC User Support Tools (HUST)*, 2022. doi: 10.1109/hust56722.2022.00010.
- [9] Pedro Valero-Lara, Aaron Young, Thomas Naughton III, Christian Engelmann, and Al Geist. ChatMPI: LLM-driven MPI code generation for HPC workloads. In *Supercomputing Asia*, 2026. doi: 10.1145/3773656.3773659.
- [10] Mingkai Zheng, Fangru Linghu, Sikan Li, Marty Charles Kandes, and Niall Gaffney. A comparative study on LLM-enabled HPC user support. In *Practice and Experience in Advanced Research Computing (PEARC)*, 2026. doi: 10.1145/3785462.3815796.